

AI Engineer Interview Questions

Real Questions from Big Tech — Sourced & Referenced

Google • Meta • Amazon • OpenAI • Microsoft • NVIDIA

Publisher

AI Engineering Insider

Sources

Glassdoor • GitHub • IGotAnOffer
Reddit • PromptLayer • InterviewQuery

Questions

52 real, sourced questions

Edition

2025–2026

Topics Covered

LLMs, RAG, Agents,
System Design, Coding, Behavioral

Companies

6 top-tier companies

All questions sourced from public interview reports, Glassdoor reviews, GitHub repositories, and expert interview guides. Direct URL references provided for every section.

Contents

1	Overview and Source Map	3
2	LLM Fundamentals and Architecture	3
3	Retrieval-Augmented Generation (RAG)	4
4	AI Agents and Multi-Agent Systems	6
5	ML System Design	6
6	Fine-Tuning, RLHF, and Alignment	7
7	LLM Engineering and Production Failures	8
8	Coding and Data Structures	8
9	ML Theory and Statistics	9
10	Behavioral and Culture Fit	10
11	NVIDIA and GPU-Specific Questions	10
12	Quick-Reference Cheatsheet	11
13	All Source References	11

1 Overview and Source Map

This document compiles 52 real AI Engineer interview questions reported by candidates at major tech companies. Each question includes company tags, difficulty rating, and a verifiable source URL.

Company	Focus Areas	Difficulty
GOOGLE	ML system design, coding (LC Hard), Gemini-era AI	5/5
META	Recommendation systems, ranking, multi-modal AI	4/5
AMAZON	Applied ML, Bedrock/RAG, SageMaker, Leadership Principles	3/5
OPENAI	LLM internals, alignment, safety, agentic systems	5/5
MSFT	Azure AI, Copilot, RAG pipelines, LangChain	4/5
NVIDIA	GPU architecture, inference optimization, CUDA	5/5

Reference Sources

- ▷ [Glassdoor — Google AI Engineer Interview Questions](#) (8 reviews, 2025)
- ▷ [Glassdoor — Google MLE Interview Questions](#) (42 reviews, 2025)
- ▷ [Glassdoor — Amazon ML Engineer Interview Questions](#) (31 reviews)
- ▷ [Glassdoor — OpenAI Software Engineer Interview](#) (38 reviews, Feb 2026)
- ▷ [GitHub: amitshekhariitbhu/ai-engineering-interview-questions](#)
- ▷ [GitHub: KalyanKS-NLP/RAG-Interview-Questions-and-Answers-Hub](#) (100+ Q&As)
- ▷ [GitHub: llmgenai/LLMInterviewQuestions](#) (Google, NVIDIA, Meta, Microsoft)
- ▷ [IGotAnOffer — 40+ AI Engineer Interview Questions](#) (2025)
- ▷ [IGotAnOffer — Google MLE Interview](#) (Glassdoor + Blind sourced)
- ▷ [PromptLayer — Agentic System Design Interview](#) (July 2025)
- ▷ [LockedIn AI — LLM/AI Engineer Interview Questions Silicon Valley](#) (Jan 2026)
- ▷ [GitHub: alexeygrigorev/ai-engineering-field-guide](#) (Q4 2025/Q1 2026)

2 LLM Fundamentals and Architecture

These questions test deep understanding of transformer internals, tokenization, attention mechanisms, and generation strategies. Reported across Google, Meta, OpenAI, and NVIDIA interviews on Glassdoor and GitHub repositories.

▷ [amitshekhariitbhu/ai-engineering-interview-questions](#) ▷ [Devinterview-io/llms-interview-questions](#)

Q1 — Tokenization and Subword Encoding

GOOGLE

META

MEDIUM

Explain Byte Pair Encoding (BPE). How does it differ from WordPiece and SentencePiece? What are the failure modes when a tokenizer splits important domain-specific terms into meaningless subword pieces, and how do you fix it?

Follow-ups: How does vocabulary size impact inference speed and memory? Why does GPT-4 use Tiktoken while LLaMA-3 uses a different vocabulary scheme?

Q2 — Attention Mechanism and KV Cache

GOOGLE

OPENAI

HARD

Explain the Query (Q), Key (K), and Value (V) matrices in multi-head attention. Your Transformer's KV cache grows too large during long-sequence generation. How do you manage memory without degrading response quality?

Follow-ups: Explain Grouped Query Attention (GQA) vs Multi-Head Latent Attention (MLA). Why did Mistral adopt sliding window attention?

Q3 — Positional Encoding and RoPE

META

NVIDIA

HARD

What is positional encoding and why is it needed in Transformers? How does Rotary Position Embedding (RoPE) work, and why is it preferred over learned positional embeddings for extending context length?

Follow-ups: How does YaRN extend RoPE for 128K+ context? What is the “lost in the middle” problem and how does it affect production RAG systems?

Q4 — Sampling Strategies

GOOGLE

OPENAI

MEDIUM

Explain Top-p (nucleus) sampling and Top-k sampling. How do they differ? What are logits, and how are they used in text generation? When would you prefer temperature scaling over beam search?

Q5 — Speculative Decoding

OPENAI

META

HARD

Your LLM inference latency is too high for real-time chat. Explain speculative decoding. How does a smaller “draft” model predict tokens that are verified in parallel by the target LLM, and what is the acceptance rate trade-off?

Q6 — Autoregressive vs. Masked Language Modeling

GOOGLE

META

MEDIUM

Explain the difference between autoregressive (causal) and masked language modeling. Why do decoder-only models dominate production LLM deployments? What is the role of the LM head in converting hidden states to vocabulary logits?

Q7 — Foundation Models

AMAZON

EASY

What are foundation models, and how have they changed AI engineering? What is the difference between open-source and closed-source LLMs? Give concrete trade-offs for a production use case (e.g., latency, cost, data privacy, customizability).

Expert Tip

At Google, interviewers often follow up transformer questions with a system design prompt: “How would you scale this for 350 million daily users?” Always connect theory to production constraints (latency, cost, throughput). Source: [IGotAnOffer — Google MLE Guide](#)

3 Retrieval-Augmented Generation (RAG)

RAG is the most heavily tested topic across all AI engineer interviews in 2025. Expect both architecture design questions and failure-mode debugging scenarios.

▷ [KalyanKS-NLP/RAG-Interview-Questions-and-Answers-Hub](#) ▷ [llmgenai/LLMInterviewQuestions](#) ▷ [IGotAnOffer](#)

Q8 — RAG Architecture

AMAZON

MSFT

MEDIUM

What is a RAG pipeline and how would you design it using Python code? Explain R (Retrieve), A (Augment), and G (Generate) step by step. What popular frameworks would you use?

Expected structure: Document ingestion → Chunking → Embedding → Vector store → Retrieval → Context injection → LLM generation.

Q9 — Chunking Strategies

GOOGLE

META

HARD

What are chunking strategies, and how do you choose the right chunk size? Compare fixed-size chunking, semantic chunking, and recursive chunking. What are the consequences of chunks that are too large vs. too small?

Follow-up: What is parent-child chunking, and how does it improve retrieval precision?

Q10 — Retrieval Quality Debugging

AMAZON

MSFT

HARD

A client's RAG system is not giving accurate results. Upon investigation, the retrieval component is inaccurate. What concrete steps do you take to improve retrieval quality?

Expected answer covers: Re-ranking, hybrid search (BM25 + dense), embedding fine-tuning, query rewriting, HyDE, chunk overlap tuning. Source: [llmgenai/LLMInterviewQuestions](#)

Q11 — Re-ranking

GOOGLE

OPENAI

HARD

A RAG system retrieves 20 candidate chunks but can only fit 5 in the LLM's context window. Without re-ranking, how does this affect response quality? What specific problems would a cross-encoder re-ranker solve that a bi-encoder cannot?

Follow-up: Describe the vector pre-computation strategy in a bi-encoder + cross-encoder pipeline. Why can't cross-encoders pre-compute representations like bi-encoders? Source: [KalyanKS-NLP/RAG-Interview-Questions-and-Answers-Hub](#)

Q12 — Hallucination and Faithfulness

META

AMAZON

HARD

Your RAG system is hallucinating despite having the right context in the retrieved chunks. How do you fix it? Explain how hallucinations specifically impact the Faithfulness metric and what techniques improve Faithfulness scores.

Follow-up: How does Response Relevancy differ from Context Relevancy, and why do you need both metrics to properly evaluate a RAG system?

Q13 — Scale and Access Control

AMAZON

MSFT

MEDIUM

How do you scale a RAG system to millions of documents? Your RAG system needs per-user access control on internal documents. How do you implement it without exposing sensitive data across user boundaries?

Q14 — Query Transformation

GOOGLE

META

HARD

What is query transformation in RAG? Explain HyDE (Hypothetical Document Embeddings), query decomposition, and step-back prompting. When would you use each technique?

Q15 — RAG vs. Fine-Tuning

OPENAI

HARD

Compare RAG vs. fine-tuning. When would you use each? What are the cost, latency, and freshness trade-offs? Can you combine both approaches?

Note: This appears frequently in OpenAI and Anthropic interviews. Interviewers expect a nuanced answer covering domain adaptation vs. knowledge injection. Source: [IGotAnOffer — AI Engineer Interview Guide](#)

Q16 — PDF Parsing and Multimodal Documents

AMAZON

MEDIUM

Your RAG system struggles with PDF documents containing tables and complex layouts. How do you fix PDF parsing? What tools or models would you use to extract structured content faithfully?

4 AI Agents and Multi-Agent Systems

Agentic system design is the fastest-growing interview topic in 2025. Questions range from ReAct architecture to multi-agent coordination and tool design.

▷ [amitshekhariitbhu/ai-engineering-interview-questions](#) ▷ [PromptLayer — Agentic System Design Interview \(July 2025\)](#)

Q17 — Agent Architecture

OPENAI

GOOGLE

HARD

What is an AI agent, and how does it differ from a simple LLM call? Explain the ReAct (Reasoning + Acting) agent architecture. What are the different types of agent memory: short-term, long-term, and episodic?

Q18 — Multi-Agent System Design (Real PromptLayer Q)

META

OPENAI

HARD

1. Design an agent that analyzes customer support tickets and automatically drafts responses, escalating complex issues to humans.
2. Create a system where agents collaborate to produce a research report with citations.
3. Build an agent that reads code, understands its purpose, and suggests improvements.

Expected: Flow chart of steps, model selection for each sub-task, draft prompt design. Source: [PromptLayer Blog \(July 2025\)](#)

Q19 — Tool Design

GOOGLE

MSFT

HARD

How do you design and define tools for an AI agent? What is the difference between single-agent and multi-agent systems? Implement a function to batch API calls to an LLM efficiently.

Q20 — Agentic RAG

OPENAI

HARD

What is Agentic RAG, and how does it differ from static RAG? How does an agentic system decide *when* to search, *which* tool to use, and *how* to aggregate multi-hop information?

Q21 — Prompt Injection Defense

META

AMAZON

HARD

Your LLM agent is vulnerable to prompt injection that reveals the system prompt. Your chatbot's proprietary business logic is being leaked by users. How do you defend against these attacks in production?

5 ML System Design

System design rounds typically last 45–60 minutes. These real questions come from Google and Meta Glassdoor/Blind reports. Interviewers expect trade-off discussions on latency, throughput, cost, and accuracy.

▷ [IGotAnOffer — Google MLE \(Glassdoor + Blind sourced\)](#) ▷ [Glassdoor — Google MLE \(42 reviews\)](#)

Q22 — Content Moderation System (Real Glassdoor Q)

GOOGLE

HARD

How would you build, train, and deploy a system that detects if multimedia and/or ad content violates terms or contains offensive materials?

Follow-ups: How do you handle multilingual content? How do you update the model as new violation categories emerge? What latency SLAs would you set?

Q23 — Autocomplete and Spell Check (Real Glassdoor Q)

GOOGLE

HARD

Design autocomplete and/or spell check on a mobile device. Then design autocomplete and/or automatic responses for email at Google scale.

Source: Reported by candidates on Glassdoor and Blind for Google MLE roles. [IGotAnOffer](#) — [Google MLE](#)

Q24 — YouTube Recommendation System

GOOGLE

META

HARD

Design the YouTube recommendation system. How would you handle the cold-start problem? How do you balance exploration vs. exploitation? What real-time and batch components exist in the architecture?

Also asked: Design a personalized content feed (Meta News Feed). Design a machine learning system for YouTube Video Search.

Q25 — Ranking and Recommendation (Real Glassdoor Q)

META

HARD

If you were creating an algorithm for a Quora-like Q&A system with the objective of ensuring users find the most pertinent answers quickly, which evaluation metric would you choose? If you have search results from multiple retrieval methods, how would you merge and homogenize rankings into a single result set?

Source: [llmgenai/LLMInterviewQuestions](#)

Q26 — AI Pipeline in Cloud

AMAZON

MSFT

MEDIUM

Describe any AI pipelines you have implemented in a cloud provider (AWS, Azure, GCP). What is fine-tuning? How would you fine-tune a model that uses hosted LLM APIs? What cost controls and monitoring would you implement?

Source: [IGotAnOffer](#) — [AI Engineer Interview Guide](#)

6 Fine-Tuning, RLHF, and Alignment

▷ [amitshekhariitbhu/ai-engineering-interview-questions](#) ▷ [LockedIn AI — Silicon Valley LLM Interviews \(Jan 2026\)](#)

Q27 — RLHF and Alignment Tax

OPENAI

META

HARD

After RLHF alignment, your LLM became safer but lost capability on hard tasks. How do you manage the alignment tax? Your RLHF-trained LLM is gaming the reward model instead of being genuinely helpful. How do you fix reward hacking?

Q28 — LoRA and PEFT Methods

META

OPENAI

HARD

What is Low-Rank Adaptation (LoRA), and why is it preferred over full fine-tuning? Explain QLoRA for memory-efficient fine-tuning on consumer GPUs. When would you choose PEFT over full fine-tuning?

Q29 — Multilingual and Cross-Lingual Transfer

GOOGLE

HARD

Your AI system works in English but fails for other languages. How do you add multilingual support? Your zero-shot cross-lingual transfer from English fails on other languages. How do you fix it?

Source: [amitshekhariitbhu/ai-engineering-interview-questions](#)

Q30 — Embedding Fine-Tuning

AMAZON

MSFT

MEDIUM

You are working with an OpenAI embedding model and benchmarking accuracy is low. How would you further improve the accuracy of the embedding search model? Walk through the steps of improving a sentence transformer used for embedding.

Source: [llmgenai/LLMInterviewQuestions](#)

7 LLM Engineering and Production Failures

These “failure mode” questions are increasingly common in 2025. Interviewers present a broken system and ask how you debug and fix it under production constraints. Extensively reported across GitHub and LockedIn AI (Jan 2026).

▷ [amitshekhariitbhu/ai-engineering-interview-questions](#) ▷ [LockedIn AI \(Jan 2026\)](#)

Q31 — Structured Output Compliance

OPENAI

AMAZON

MEDIUM

Your LLM keeps ignoring your instructions. How do you make it follow structured output formats? What is the difference between zero-shot prompting and few-shot prompting with format examples?

Q32 — Context Window Management

META

GOOGLE

MEDIUM

Your LLM-powered tool hits the context window limit on long documents. How do you handle it? Your chatbot loses context after 10 turns. How do you maintain long conversation context effectively?

Q33 — Hallucination and Uncertainty

OPENAI

MEDIUM

Your LLM does not admit when it does not know the answer. How do you make it say “I don’t know”? Your QA system always generates an answer even when no answer exists in the context. How do you detect unanswerable questions?

Q34 — Cost and Latency Optimization (Real SV Interview Q)

AMAZON

MSFT

MEDIUM

With GPT-4 Turbo processing costing approximately \$13,000 per day for 2 billion tokens, discuss strategies like caching, model tiering, and prompt optimization to reduce cost. How do you decide when to use a smaller model vs. a frontier model?

Source: [LockedIn AI Silicon Valley Guide \(January 2026\)](#)

Q35 — Prompt Sensitivity

GOOGLE

META

HARD

Your LLM classification system is too sensitive to prompt wording changes. How do you reduce prompt sensitivity? Explain how chain-of-thought prompting improves reasoning, and when it fails to improve LLM accuracy.

Q36 — Data Leakage and Privacy

AMAZON

MEDIUM

Your LLM memorized proprietary training data and leaks it in responses. How do you prevent this? What is differential privacy, and how can it be applied during LLM training?

8 Coding and Data Structures

Coding rounds at AI engineer interviews mirror SWE rounds but include AI-specific practical tasks. Google focuses on LeetCode Hard; OpenAI prefers real-world data structure problems over abstract puzzles.

▷ [IGotAnOffer — Google MLE Coding \(Glassdoor data\)](#) ▷ [Glassdoor — OpenAI SWE \(38 reviews\)](#)

Q37 — Binary Tree Maximum Path Sum (Real Google Q)

GOOGLE

HARD

Given a binary tree, find the maximum path sum.

Source: Reported by multiple Google MLE candidates on Glassdoor. [IGotAnOffer](#) — [Google MLE Interview](#)

Q38 — Longest Substring Without Repeating Characters (Real OpenAI Q)

OPENAI

MEDIUM

Find the length of the longest substring without repeating characters.

Source: Reported by an OpenAI SWE candidate on Glassdoor, Feb 2026. [Glassdoor](#) — [OpenAI](#)

Q39 — Versioned Data Store (Real OpenAI Q)

OPENAI

MSFT

HARD

Implement a time-based data structure or versioned key-value store.

Note: OpenAI prefers practical engineering tasks that reflect real-world scenarios over abstract competitive programming puzzles. Source: [LockedIn AI](#) (Jan 2026)

Q40 — Numerical Precision

GOOGLE

HARD

What does it mean for an algorithm to be numerically unstable? What are the trade-offs between FP32, FP16, and BF16 precision when training large-scale models on NVIDIA GPUs?

Follow-up: When dealing with massive datasets, how do you decide between sparse matrix vs. dense representations, and how does this affect memory bandwidth? Source: [IGotAnOffer](#)

Q41 — Vector Search and Indexing

META

AMAZON

MEDIUM

You need to find similar customer reviews and perfect accuracy is the priority. Which search strategy do you choose and why? Explain vector search strategies like clustering and Locality-Sensitive Hashing. Compare different vector index types and when you would use each.

Source: [llmgenai/LLMInterviewQuestions](#)

9 ML Theory and Statistics

▷ [amitshekhariitbhu/machine-learning-interview-questions](#) ▷ [Devinterview-io/llms-interview-questions](#)

Q42 — Overfitting and Regularization

GOOGLE

AMAZON

MEDIUM

Explain overfitting and underfitting. How can you prevent them? What is regularization? Explain L1 (Lasso) and L2 (Ridge) regularization and their geometric interpretations.

Q43 — Self-Attention Complexity

META

GOOGLE

MEDIUM

Explain the disadvantages of the self-attention mechanism (quadratic complexity w.r.t. sequence length). How can you overcome it? Name at least three efficient attention variants and their trade-offs.

Q44 — Evaluation Metrics for Information Retrieval

AMAZON

MEDIUM

Compare information retrieval metrics: Precision@K, Recall@K, MRR, NDCG, and MAP. Which metric should you use when? If you have a recommendation system, which metric would you use to evaluate it?

Source: [llmgenai/LLMInterviewQuestions](#)

Q45 — Multimodal Architecture

GOOGLE

META

HARD

Transformers work on text, so can they also understand images? Explain how Vision Transformers (ViT) work. How do modern multimodal models like GPT-4o handle interleaved text, image, and voice inputs without

requiring separate specialized encoders?

10 Behavioral and Culture Fit

At Meta, the behavioral interview can constitute nearly half of the overall evaluation score. Failing the behavioral round can mean falling short overall, even with strong technical performance.

▷ [LockedIn AI \(Jan 2026\)](#) ▷ [PromptLayer — AI Engineer Evaluation \(July 2025\)](#)

Q46 — Model Opinions and Current Knowledge

OPENAI

GOOGLE

MEDIUM

Which model provider do you lean on for creative writing: Claude, OpenAI, or Gemini? What AI coding assistants do you use? Can you compare Cursor, Windsurf, and Claude Code? How do open-source LLMs stack up against proprietary ones?

Tip: Interviewers look for opinionated candidates with well-reasoned trade-offs. Source: [PromptLayer Blog \(July 2025\)](#)

Q47 — Project Impact

GOOGLE

EASY

Describe a machine learning project you worked on and its impact. How would you explain machine learning to someone who doesn't understand it?

Q48 — Situational Problem-Solving

AMAZON

MEDIUM

What would you do if a dataset became corrupted mid-training? How would you handle GPU saturation during a critical project? How would you mitigate a prompt injection attack discovered in a live production system?

Source: [LockedIn AI Silicon Valley Guide \(Jan 2026\)](#)

Q49 — AI Ethics and Safety

META

OPENAI

MEDIUM

How do you ensure responsible AI? How do you integrate fairness, transparency, and privacy principles into every stage of model design and deployment? What is your approach to hallucination mitigation and bias detection in production systems?

11 NVIDIA and GPU-Specific Questions

▷ [Devinterview-io/llms-interview-questions](#) ▷ [LockedIn AI \(Jan 2026\)](#)

Q50 — Training Parallelism

NVIDIA

GOOGLE

HARD

Explain 3D Parallelism: Data Parallelism, Tensor Parallelism, and Pipeline Parallelism. How do frameworks like Megatron-LM and PyTorch FSDP2 implement these? When would you choose each strategy for a 70B parameter model?

Q51 — Quantization Techniques

NVIDIA

HARD

Explain GPTQ, AWQ, and QuIP# quantization. When does training in BFloat16 or FP8 using Transformer Engine (TE) maximize throughput on H100/B200 clusters? What are the accuracy trade-offs?

Q52 — FlashAttention and Memory Efficiency

NVIDIA

GOOGLE

HARD

What is FlashAttention, and how does it reduce memory bandwidth bottlenecks in transformer training? Your Transformer runs out of memory on long documents due to quadratic self-attention. How do you scale it?

12 Quick-Reference Cheatsheet

LLM Concepts Checklist

- BPE, WordPiece, SentencePiece
- Multi-head attention, GQA, MLA
- KV cache and memory management
- RoPE, YaRN positional encoding
- Top-k, top-p, temperature sampling
- Speculative decoding
- RLHF, DPO, ORPO alignment
- LoRA, QLoRA, DoRA PEFT methods
- Causal vs. masked LM pre-training
- FlashAttention, Paged Attention

RAG and Agent Checklist

- Chunking: fixed, semantic, recursive
- Bi-encoder vs. cross-encoder
- Hybrid search: BM25 + dense
- Re-ranking strategies
- HyDE, query decomposition
- RAGAS evaluation metrics
- ReAct agent pattern
- Tool design for agents
- Multi-agent coordination
- Prompt injection defense

Expert Tip

Top Tips from 200+ Coached Candidates (Viral, Ex-Meta/LinkedIn Engineering Leader):

- “Companies are not hiring for titles — they want to know if you can build reliable AI systems.”
- “If you can only do ML *or* only system design, you’re already behind. The industry is moving to AI system builders.”
- “Think like a product + system owner, not just someone calling an LLM API.”
- Always connect theory to production: latency, cost, throughput, and failure modes.

Source: [IGotAnOffer — AI Engineer Interview Guide \(2025\)](#)

13 All Source References

1. **Glassdoor — Google AI Engineer Interview Questions** (8 reviews, Oct 2025)
https://www.glassdoor.com/Interview/Google-AI-Engineer-Interview-Questions-EI_IE9079.0,6_K07,18.htm
2. **Glassdoor — Google Machine Learning Engineer Interview** (42 reviews, Aug 2025)
https://www.glassdoor.com/Interview/Google-Machine-Learning-Engineer-Interview-Questions-EI_IE9079.0,6_K07,32.htm
3. **Glassdoor — Amazon Machine Learning Engineer Interview** (31 reviews)
https://www.glassdoor.com/Interview/Amazon-Machine-Learning-Engineer-Interview-Questions-EI_IE6036.0,6_K07,32.htm
4. **Glassdoor — OpenAI Software Engineer Interview** (38 reviews, Feb 2026)
https://www.glassdoor.com/Interview/OpenAI-Software-Engineer-Interview-Questions-EI_IE2210885.0,6_K07,24.htm
5. **Glassdoor — Meta Machine Learning and Computer Vision Engineer** (2 reviews)
https://www.glassdoor.com/Interview/Meta-Machine-Learning-and-Computer-Vision-Engineer-Interview-Questions-EI_IE40772.0,4_K05,50.htm
6. **GitHub: amitshekhariitbhu/ai-engineering-interview-questions** (IIT Bombay alumnus)
<https://github.com/amitshekhariitbhu/ai-engineering-interview-questions>
7. **GitHub: KalyanKS-NLP/RAG-Interview-Questions-and-Answers-Hub** (100+ RAG Q&As)
<https://github.com/KalyanKS-NLP/RAG-Interview-Questions-and-Answers-Hub>
8. **GitHub: llmgenai/LLMInterviewQuestions** (Google, NVIDIA, Meta, Microsoft, Fortune 500)

<https://github.com/llmgenai/LLMInterviewQuestions>

9. **GitHub: Devinterview-io/llms-interview-questions** (LLM Interview Q&A 2026)
<https://github.com/Devinterview-io/llms-interview-questions>
10. **GitHub: a-tabaza/genai_interview_questions** (Theory-grounded GenAI questions)
https://github.com/a-tabaza/genai_interview_questions
11. **GitHub: alexeygrigorev/ai-engineering-field-guide** (Q4 2025/Q1 2026 hiring research)
<https://github.com/alexeygrigorev/ai-engineering-field-guide>
12. **IGotAnOffer — 40+ AI Engineer Interview Questions** (Expert: Viral, Ex-Meta/LinkedIn)
<https://igotanooffer.com/en/advice/ai-engineer-interview>
13. **IGotAnOffer — Google Machine Learning Engineer Interview** (Glassdoor + Blind sourced)
<https://igotanooffer.com/blogs/tech/google-machine-learning-engineer-interview>
14. **PromptLayer Blog — Agentic System Design Interview** (July 2025)
<https://blog.promptlayer.com/the-agentic-system-design-interview-how-to-evaluate-ai-engineers/>
15. **LockedIn AI — LLM & AI Engineer Interview Questions Silicon Valley** (January 2026)
<https://www.lockedinai.com/blog/llm-ai-engineer-interview-questions-silicon-valley>
16. **InterviewQuery — Google Machine Learning Engineer Interview Guide**
<https://www.interviewquery.com/interview-guides/google-machine-learning-interview-questions>
17. **InterviewQuery — Google AI Engineer Interview Guide (2026)**
<https://www.interviewquery.com/prep-guides/google-ai-engineer>

All questions are sourced from public interview reports and community-maintained repositories. This document is for educational purposes only.

AI Engineering Insider | aiengineeringinsider.substack.com | beacons.ai/aiengineeringinsider | linkedin.com/in/lamhotsiagian